

Assignment 3

ANS-1

- a) Not Stationary as mean is not constant (it is increasing).
- b) Stationary as all conditions of stationarity are satisfied.
- c) Not Stationary as variance is different in the middle portion.
- d) Not Stationary as the data is periodic.
- e) Not Stationary as mean is not constant (is is decreasing).
- f) Not Stationary as there is a peak at the beginning.
- g) Stationary as the data is periodic with irregular intervals.
- h) Not Stationary as the data is periodic.
- i) Not Stationary as mean is increasing.

ANS-2

```
library('forecast')
```

```
## Registered S3 method overwritten by 'xts':  
##   method      from  
##   as.zoo.xts zoo
```

```
## Registered S3 method overwritten by 'quantmod':  
##   method      from  
##   as.zoo.data.frame zoo
```

```
## Registered S3 methods overwritten by 'forecast':  
##   method      from  
##   fitted.fracdiff fracdiff  
##   residuals.fracdiff fracdiff
```

```
library('tseries')
```


```
amazon=read.csv('amazon.csv')  
amazon=amazon[which(amazon$date=='26/10/2016'):nrow(amazon),'close']  
amazon=ts(amazon,start=c(2016,10,26),frequency=365)  
  
google=read.csv('google.csv')  
google=google[which(google$date=='26/10/2016'):nrow(google),'close']  
google=ts(google,start=c(2016,10,26),frequency=365)  
  
micro=read.csv('microsoft.csv')
```

```
micro[micro$date=='26/10/2016']:nrow(micro), 'close']
micro=ts(micro, start=c(2016, 10, 26), frequency=365)

plot(amazon, type = 'l', ylab='Price per share (in USD)', main = 'Amazon stock price')
```




```
plot(google, type = 'l', ylab='Price per share (in USD)', main = 'Google stock price')
```



```
plot(micro,type = 'l',ylab='Price per share (in USD)',main = 'Microsoft stock price')
```

Microsoft stock price




- a) All the stocks are performing very well as they are showing an upward trend. Microsoft's stocks are very stable, whereas Google's stock shows a small periodic decline and generally has higher variance.

Amazon's stock value is fluctuating around a constant value for the past 1 year and hasn't seen growth recently.


- b)

```
acf_amazon=acf(amazon,lag.max=100, plot=FALSE)
pacf_amazon=pacf(amazon,lag.max=100, plot=FALSE)
par(mfrow=c(2,1))
plot(acf_amazon)
plot(pacf_amazon)
```

Series amazon




Series amazon




```
acf_google=acf(google,lag.max=100, plot=FALSE)
pacf_google=pacf(google,lag.max=100, plot=FALSE)
par(mfrow=c(2,1))
plot(acf_google)
plot(pacf_google)
```

Series google




Series google




```
acf_micro=acf(micro,lag.max=100, plot=FALSE)
pacf_micro=pacf(micro,lag.max=100, plot=FALSE)
par(mfrow=c(2,1))
plot(acf_micro)
plot(pacf_micro)
```

Series micro



Series micro




In all of the above cases we see that auto correlation is decreasing with increasing lags. In partial auto-correlation we see that the autocorrelation is high at lag=1. This implies that value of y at t depends only on the value at $t-1$. $q=1$ for our AR model.

c)


```
amazon_stat=diff(amazon,differences = 1)
google_stat=diff(google,differences = 1)
micro_stat=diff(micro,differences = 1)
plot(amazon_stat, main='Amazon')
```

Amazon




```
plot(google_stat, main='Google')
```


Google



```
plot(micro_stat, main='Microsoft')
```



All of the above can be made approximately stationary by differencing the data by lag=1.

- d) Learn only from first 600 data points and then use it to predict future samples. Finally, calculate mae error.

```
# Amazon prices
amazon_fit=arima(amazon[1:600],order=c(1,1,1))
forecast_amazon=forecast(amazon_fit,h=155)
sprintf('MAE error on future samples of Amazon: %s',sum(abs(forecast_amazon[["mean"]]-amazon[601:755])))

## [1] "MAE error on future samples of Amazon: 17409.0814325496"

google_fit=arima(google[1:600],order=c(1,1,1))
forecast_google=forecast(google_fit,h=155)
sprintf('MAE error on future samples of Google: %s',sum(abs(forecast_google[["mean"]]-google[601:755])))

## [1] "MAE error on future samples of Google: 6787.32233402785"

micro_fit=arima(micro[1:600],order=c(1,1,1))
forecast_micro=forecast(micro_fit,h=155)
sprintf('MAE error on future samples of Microsoft: %s',sum(abs(forecast_micro[["mean"]]-micro[601:755])))

## [1] "MAE error on future samples of Microsoft: 2368.64804756622"
```

ANS-3

```
# load the data
library(mlbench)
data("BreastCancer")
cancer=data.frame(data.matrix(BreastCancer[1:nrow(BreastCancer),-1]))
cancer[is.na(cancer)]<-1      # Replace NA values with mode only bare.nuclei columns has na values

model=glm(as.factor(Class)~.,data=cancer, family='binomial')

# predicted=predict.glm(model,cancer,type='response')
# predicted=ifelse(predicted>0.5,2,1)

# [501:699,1:9]
# (predicted>0.5)==cancer[501:699,10]
# table(cancer[501:699,10],predicted>0.5)

# sum(predicted==cancer[1:699,10])

summ=summary.glm(model)
summ

##
## Call:
## glm(formula = as.factor(Class) ~ ., family = "binomial", data = cancer)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -3.3324  -0.1269  -0.0658   0.0272   2.3958
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -9.72056    1.07329  -9.057  < 2e-16 ***
## Cl.thickness    0.53489    0.13482   3.967  7.27e-05 ***
## Cell.size       0.01162    0.19327   0.060  0.95205
## Cell.shape      0.32323    0.21391   1.511  0.13077
## Marg.adhesion   0.23765    0.11672   2.036  0.04175 *
## Epith.c.size    0.05850    0.15252   0.384  0.70131
## Bare.nuclei     0.42816    0.09054   4.729  2.25e-06 ***
## Bl.cromatin     0.41228    0.15694   2.627  0.00862 **
## Normal.nucleoli 0.15826    0.10416   1.519  0.12866
## Mitoses         0.53953    0.30397   1.775  0.07591 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 900.53  on 698  degrees of freedom
## Residual deviance: 113.10  on 689  degrees of freedom
## AIC: 133.1
##
## Number of Fisher Scoring iterations: 8
```

We fit a logit link function to the data as the output variable is a binary variable and obtain the summary as above.

Estimate column tells us the coefficients of each independent variable including a bias term.

Std. Error is the standard deviation of that variable.

z-value is obtained by dividing the estimate and the std. error.

Pr(>|z|) tells the probability of obtaining that z value.

Variables with *'s at the end denote that they are statistically significant with a significance level of 0.05.

```
prob=summ[["coefficients"]][31:40]
oddsratio=(prob/(1-prob))**2
oddsratio
```

```
## [1] 1.803638e-38 5.279680e-09 3.942752e+02 2.263327e-02 1.898273e-03
## [6] 5.512622e+00 5.079552e-12 7.551718e-05 2.180144e-02 6.746989e-03
```

Variables with high odds ratio also have low value of coefficients meaning these variables are less important in predicting the output class.

More the important a variable is in predicting the class less the odds ratio it has.

```
# mean=summ[["coefficients"]][1:10]
std=summ[["coefficients"]][11:20]

low=oddsratio-(std*1.96)/(sqrt(nrow(cancer)))
high=oddsratio+(std*1.96)/(sqrt(nrow(cancer)))

low
```

```
## [1] -0.079566975 -0.009994731 394.260892682 0.006775216 -0.006754885
## [6] 5.501315250 -0.006711759 -0.011559318 0.014079600 -0.015787411
```

```
high
```

```
## [1] 7.956698e-02 9.994741e-03 3.942895e+02 3.849133e-02 1.055143e-02
## [6] 5.523930e+00 6.711759e-03 1.171035e-02 2.952328e-02 2.928139e-02
```